

Installation instructions for Hadoop & MapReduce Lab

Contents

[Contents](#)

[Downloading the VM \(includes data sets\)](#)

[Virtual machine setup](#)

[Using Oracle VirtualBox](#)

[Using VMWare](#)

[Alternative: install of single node Hadoop cluster \(no virtual machine\)](#)

Downloading the VM (includes data sets)


1. Download it from
<http://content.udacity-data.com/courses/ud617/Cloudera-Udacity-Training-VM-4.1.1.c.zip>
Warning - the zipped file size is 1.7 GB. *If you are on a Windows machine you will likely need to use WinRAR to open this .zip file because other methods fail to open the unzipped file (which exceeds the maximum specified 4GB for a .zip file).*
2. {optional} MD5sum file can be found here (in case you want to verify that the download happened properly)
<http://content.udacity-data.com/courses/ud617/Cloudera-Udacity-Training-VM-4.1.1.c.zip.md5>
3. Unzip it. Warning - the unzipped size is 4.2GB
4. MD5 hashes for files:
 - o 8a610c151d4b1ebdce11542d13dd2a53 Cloudera-Training-VM-4.1.1.c.log
 - o 6b44c965c1c6062554bf4cc12d11e87e Cloudera-Training-VM-4.1.1.c.plist
 - o 46dedeba3e0affd8311431d7e370705e Cloudera-Training-VM-4.1.1.c.vmdk
 - o d41d8cd98f00b204e9800998ecf8427e Cloudera-Training-VM-4.1.1.c.vmsd
 - o 096956c1cbabeaa652ca63a2d5e14612 Cloudera-Training-VM-4.1.1.c.vmx
 - o c9f8a375e82ef1e9d96097850e237df9 Cloudera-Training-VM-4.1.1.c.vmx.f
 - o 0d7c8becb5a515068e81bb303c794e4f nvram

Virtual machine setup

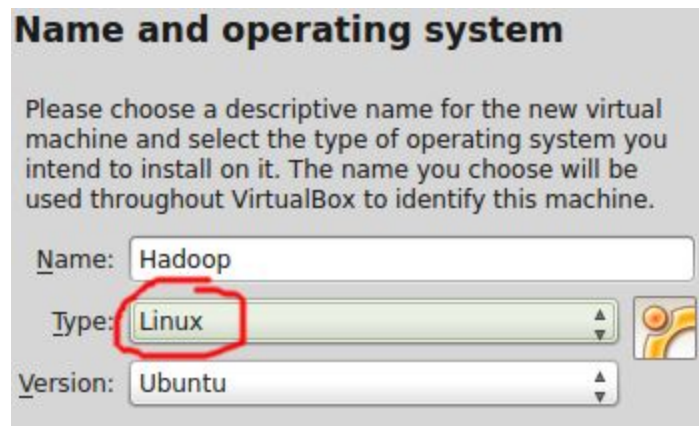
Using Oracle VirtualBox

(you can also use [brew](#) or [apt-get](#) to install vbox instead of the below)


1. Download and install VirtualBox from <https://www.virtualbox.org/wiki/Downloads>
2. Create a new Virtual machine:
 - a. Create a new virtual machine by pressing the 'New' button:



- b. Choose a name, use 'Type': 'Linux':



- c. Press Next
 - d. Select memory size for the VM.



- e. Press Next
 - f. Select 'Use an existing virtual hard drive file', click the button to browse to the directory you unzipped the provided VM image and press 'Create'.

Hard drive

If you wish you can add a virtual hard drive to the new machine. You can either create a new hard drive file or select one from the list or from another location using the folder icon.

If you need a more complex storage set-up you can skip this step and make the changes to the machine settings once the machine is created.

The recommended size of the hard drive is **8,00 GB**.

☐ Do not add a virtual hard drive

☐ Create a virtual hard drive now

☒ Use an existing virtual hard drive file

Cloudera-Training-VM-4.1.1.c.vmdk (Normal)

g. Start the VM!



Using VMWare

1. Download and install from https://my.vmware.com/web/vmware/free#desktop_end_user_computing/vmware_player/6_0
2. Create the Virtual Machine:

- a. Click on 'Open a Virtual Machine' and, when prompted, navigate to the folder you unzipped the VM, choose the file and click 'Open'.



- b. Select the machine and click 'Play virtual machine'


Cloudera-Udacity-Training- VM-4.1.1.c

State: Powered Off

OS: Other

Version: Workstation 5.x virtual machine

RAM: 1 GB



Alternative: install of single node Hadoop cluster (no virtual machine)

This can be done on any Unix machine (Mac or Linux), available on brew and apt-get. However, so extra (non-trivial) steps are involved:

<https://hadoop.apache.org/docs/stable/hadoop-project-dist/hadoop-common/SingleCluster.html>

Download datasets

Not necessary if you are using the virtual machine (they are included in the image). If not: download and unzip data sets:

- Access_log.gz
- Purchases.txt.gz

from <http://dorienherremans.com/drop/?dir=CDS/hadoop>