

intro_to_pandas

September 26, 2019

1 Intro to pandas

Learning Objectives: * Gain an introduction to the `DataFrame` and `Series` data structures of the *pandas* library * Access and manipulate data within a `DataFrame` and `Series` * Import CSV data into a *pandas* `DataFrame` * Reindex a `DataFrame` to shuffle data

pandas is a column-oriented data analysis API. It's a great tool for handling and analyzing input data, and many ML frameworks support *pandas* data structures as inputs. Although a comprehensive introduction to the *pandas* API would span many pages, the core concepts are fairly straightforward, and we'll present them below. For a more complete reference, the [pandas docs site](#) contains extensive documentation and many tutorials.

1.1 Basic Concepts

The following line imports the *pandas* API and prints the API version:

```
[1]: from __future__ import print_function

import pandas as pd
pd.__version__
```

```
[1]: '0.25.1'
```

The primary data structures in *pandas* are implemented as two classes:

- **DataFrame**, which you can imagine as a relational data table, with rows and named columns.
- **Series**, which is a single column. A `DataFrame` contains one or more `Series` and a name for each `Series`.

The data frame is a commonly used abstraction for data manipulation. Similar implementations exist in [Spark](#) and [R](#).

One way to create a `Series` is to construct a `Series` object. For example:

```
[2]: pd.Series(['San Francisco', 'San Jose', 'Sacramento'])
```

```
[2]: 0    San Francisco
     1      San Jose
     2    Sacramento
```

dtype: object

DataFrame objects can be created by passing a dict mapping string column names to their respective Series. If the Series don't match in length, missing values are filled with special NA/NaN values. Example:

```
[3]: city_names = pd.Series(['San Francisco', 'San Jose', 'Sacramento'])
      population = pd.Series([852469, 1015785, 485199])

      pd.DataFrame({ 'City name': city_names, 'Population': population })
```

```
[3]:      City name  Population
0  San Francisco    852469
1     San Jose    1015785
2  Sacramento    485199
```

But most of the time, you load an entire file into a DataFrame. The following example loads a file with California housing data. Run the following cell to load the data and create feature definitions:

```
[4]: california_housing_dataframe = pd.read_csv("https://download.mlcc.google.com/
      ↪mledu-datasets/california_housing_train.csv", sep=",")
      california_housing_dataframe.describe()
```

```
[4]:      longitude      latitude  housing_median_age  total_rooms  \
count  17000.000000  17000.000000      17000.000000  17000.000000
mean    -119.562108    35.625225        28.589353    2643.664412
std       2.005166     2.137340        12.586937    2179.947071
min     -124.350000    32.540000         1.000000     2.000000
25%     -121.790000    33.930000        18.000000   1462.000000
50%     -118.490000    34.250000        29.000000   2127.000000
75%     -118.000000    37.720000        37.000000   3151.250000
max     -114.310000    41.950000        52.000000  37937.000000

      total_bedrooms  population  households  median_income  \
count  17000.000000  17000.000000  17000.000000  17000.000000
mean     539.410824   1429.573941    501.221941     3.883578
std     421.499452   1147.852959    384.520841     1.908157
min       1.000000     3.000000     1.000000     0.499900
25%     297.000000    790.000000    282.000000     2.566375
50%     434.000000   1167.000000    409.000000     3.544600
75%     648.250000   1721.000000    605.250000     4.767000
max    6445.000000  35682.000000   6082.000000    15.000100

      median_house_value
count      17000.000000
mean     207300.912353
std     115983.764387
min      14999.000000
```

25%	119400.000000
50%	180400.000000
75%	265000.000000
max	500001.000000

The example above used `DataFrame.describe` to show interesting statistics about a `DataFrame`. Another useful function is `DataFrame.head`, which displays the first few records of a `DataFrame`:

```
[5]: california_housing_dataframe.head()
```

```
[5]:
```


	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	\
0	-114.31	34.19	15.0	5612.0	1283.0	
1	-114.47	34.40	19.0	7650.0	1901.0	
2	-114.56	33.69	17.0	720.0	174.0	
3	-114.57	33.64	14.0	1501.0	337.0	
4	-114.57	33.57	20.0	1454.0	326.0	

	population	households	median_income	median_house_value
0	1015.0	472.0	1.4936	66900.0
1	1129.0	463.0	1.8200	80100.0
2	333.0	117.0	1.6509	85700.0
3	515.0	226.0	3.1917	73400.0
4	624.0	262.0	1.9250	65500.0

Another powerful feature of *pandas* is graphing. For example, `DataFrame.hist` lets you quickly study the distribution of values in a column:

```
[7]: california_housing_dataframe.hist('housing_median_age')
```

```
[7]: array([[<matplotlib.axes._subplots.AxesSubplot object at 0x11816e390>]],
          dtype=object)
```



1.2 Accessing Data

You can access DataFrame data using familiar Python dict/list operations:

```
[11]: cities = pd.DataFrame({ 'City name': city_names, 'Population': population })
      print(type(cities['City name']))
      cities['City name']
```

```
<class 'pandas.core.series.Series'>
```

```
[11]: 0    San Francisco
      1      San Jose
      2    Sacramento
      Name: City name, dtype: object
```

```
[9]: print(type(cities['City name'][1]))
      cities['City name'][1]
```

```
<class 'str'>
```

```
[9]: 'San Jose'
```

```
[10]: print(type(cities[0:2]))
      cities[0:2]
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
[10]:      City name  Population
0  San Francisco    852469
1    San Jose      1015785
```

In addition, *pandas* provides an extremely rich API for advanced [indexing and selection](#) that is too extensive to be covered here.

1.3 Manipulating Data

You may apply Python's basic arithmetic operations to **Series**. For example:

```
[11]: population / 1000.
```

```
[11]: 0      852.469
      1     1015.785
      2      485.199
      dtype: float64
```

NumPy is a popular toolkit for scientific computing. *pandas Series* can be used as arguments to most NumPy functions:

```
[12]: import numpy as np

      np.log(population)
```

```
[12]: 0      13.655892
      1      13.831172
      2      13.092314
      dtype: float64
```

For more complex single-column transformations, you can use **Series.apply**. Like the Python [map function](#), **Series.apply** accepts as an argument a [lambda function](#), which is applied to each value.

The example below creates a new **Series** that indicates whether **population** is over one million:

```
[13]: population.apply(lambda val: val > 1000000)
```

```
[13]: 0      False
      1       True
      2      False
      dtype: bool
```

Modifying **DataFrames** is also straightforward. For example, the following code adds two **Series** to an existing **DataFrame**:

```
[14]: cities['Area square miles'] = pd.Series([46.87, 176.53, 97.92])
      cities['Population density'] = cities['Population'] / cities['Area square_
      ↪miles']
      cities
```

```
[14]:
```

	City name	Population	Area square miles	Population density
0	San Francisco	852469	46.87	18187.945381
1	San Jose	1015785	176.53	5754.177760
2	Sacramento	485199	97.92	4955.055147

2 Exercise #1

Modify the `cities` table by adding a new boolean column that is True if and only if *both* of the following are True:

- The city is named after a saint.
- The city has an area greater than 50 square miles.

Note: Boolean `Series` are combined using the bitwise, rather than the traditional boolean, operators. For example, when performing *logical and*, use `&` instead of `and`.

Hint: "San" in Spanish means "saint."

```
[0]: # Your code here
```

2.1 Indexes

Both `Series` and `DataFrame` objects also define an `index` property that assigns an identifier value to each `Series` item or `DataFrame` row.

By default, at construction, *pandas* assigns index values that reflect the ordering of the source data. Once created, the index values are stable; that is, they do not change when data is reordered.

```
[17]: city_names.index
```

```
[17]: RangeIndex(start=0, stop=3, step=1)
```

```
[18]: cities.index
```

```
[18]: RangeIndex(start=0, stop=3, step=1)
```

Call `DataFrame.reindex` to manually reorder the rows. For example, the following has the same effect as sorting by city name:

```
[19]: cities.reindex([2, 0, 1])
```

```
[19]:
```

	City name	Population	...	Population density	Is wide and has saint
name					
2	Sacramento	485199	...	4955.055147	
					False
0	San Francisco	852469	...	18187.945381	
					False
1	San Jose	1015785	...	5754.177760	
					True

[3 rows x 5 columns]

Reindexing is a great way to shuffle (randomize) a `DataFrame`. In the example below, we take the index, which is array-like, and pass it to NumPy's `random.permutation` function, which shuffles its values in place. Calling `reindex` with this shuffled array causes the `DataFrame` rows to be shuffled in the same way. Try running the following cell multiple times!

```
[20]: cities.reindex(np.random.permutation(cities.index))
```

```
[20]:
```

	City name	Population	...	Population density	Is wide and has saint
name					
0	San Francisco	852469	...	18187.945381	
					False
1	San Jose	1015785	...	5754.177760	
					True
2	Sacramento	485199	...	4955.055147	
					False

[3 rows x 5 columns]

For more information, see the [Index documentation](#).