



SINGAPORE UNIVERSITY OF  
TECHNOLOGY AND DESIGN

Established in collaboration with MIT

# Recurrent Neural Networks (Part I)

---

*SOUJANYA PORIA*

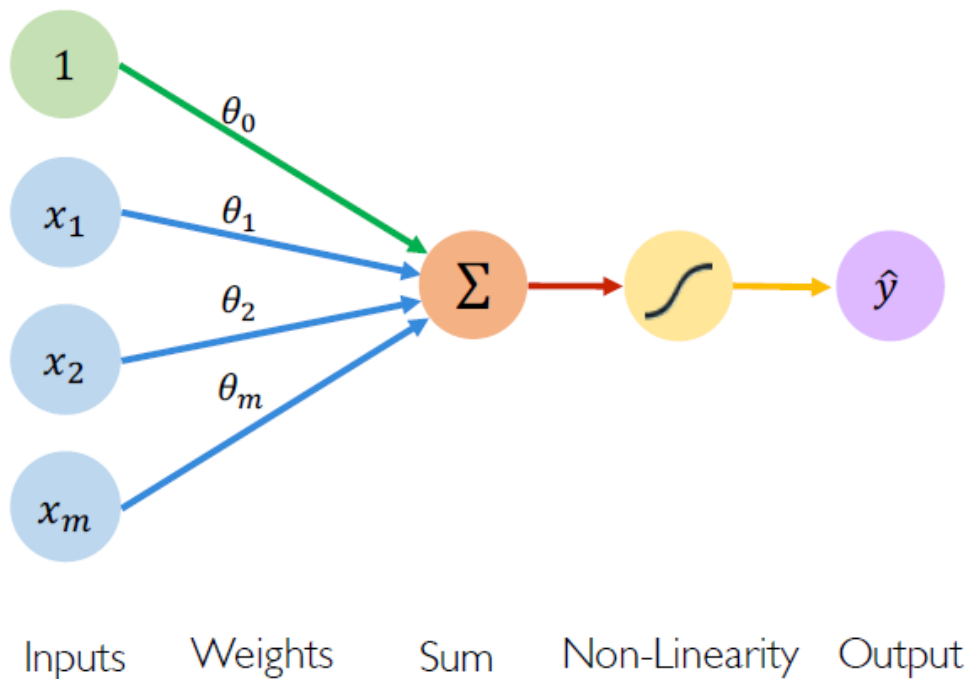
50.038 Computational data science

# Objectives

---

- Understand the limitations of feed-forward neural networks
- Understand how a Recurrent Neural Network (RNN) works
- Able to list the types of RNNs in terms of input/outputs

# Recap on Perceptron



Output

Linear combination of inputs

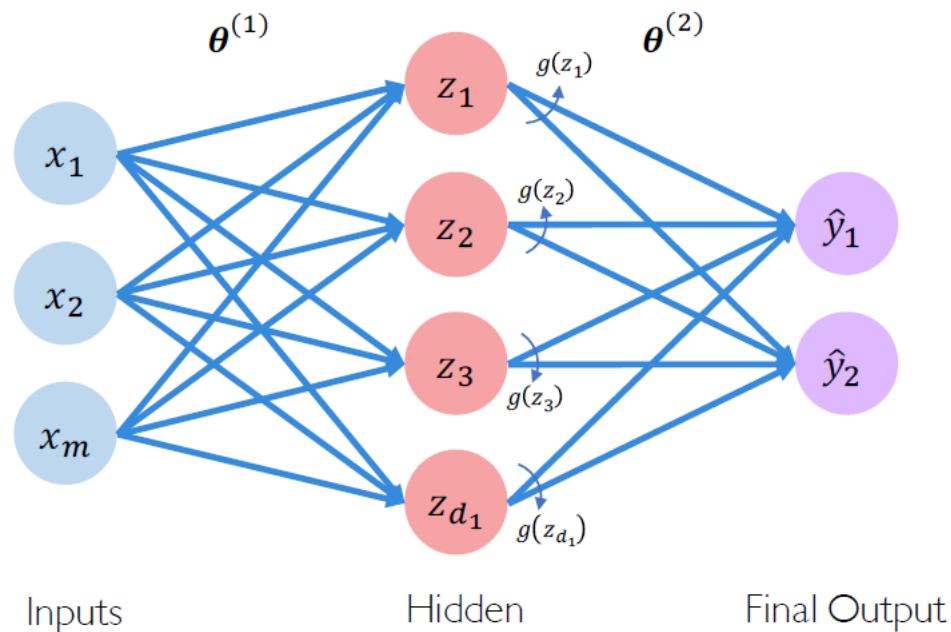
$$\hat{y} = g \left( \theta_0 + \sum_{i=1}^m x_i \theta_i \right)$$

Non-linear activation function

Bias

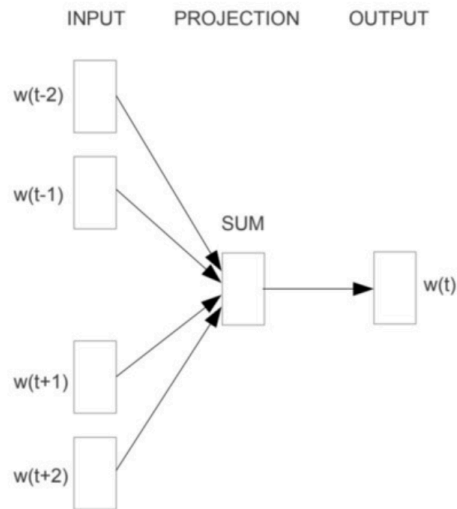
# Recap on MLP

- Multiple perceptrons, aka a Multi-layer Perceptron (MLP)

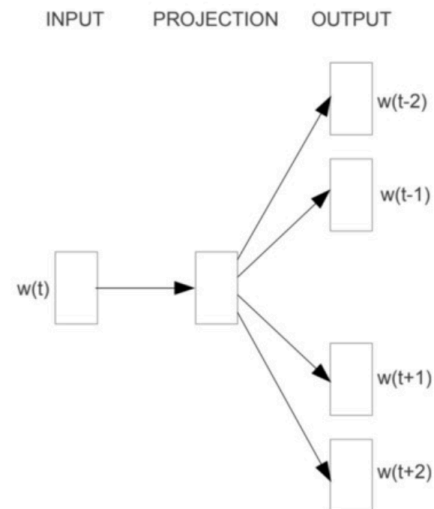


# Recap on Word2Vec

- Continuous Bag of Words (CBOW) and Skip-gram



**CBOW**



**Skip-gram**

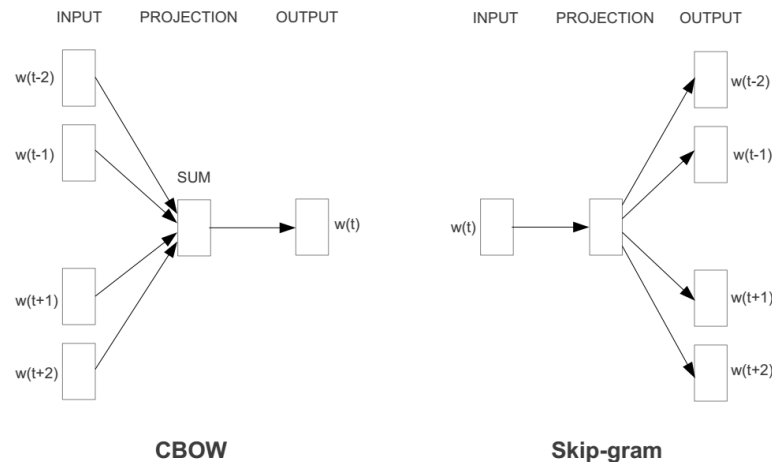
# word2vec approach to represent the meaning of word

---

- Represent each word with a low-dimensional vector
- Word similarity = vector similarity
- Key idea: Predict surrounding words of every word
- Faster and can easily incorporate a new sentence/document or add a word to the vocabulary


# Represent the meaning of word – word2vec

- 2 basic neural network models:
  - Continuous Bag of Word (CBOW): use a window of word to predict the middle word
  - Skip-gram (SG): use a word to predict the surrounding ones in window.




# Word2vec – Continuous Bag of Word


- E.g. “The cat sat on floor”
  - Window size = 2








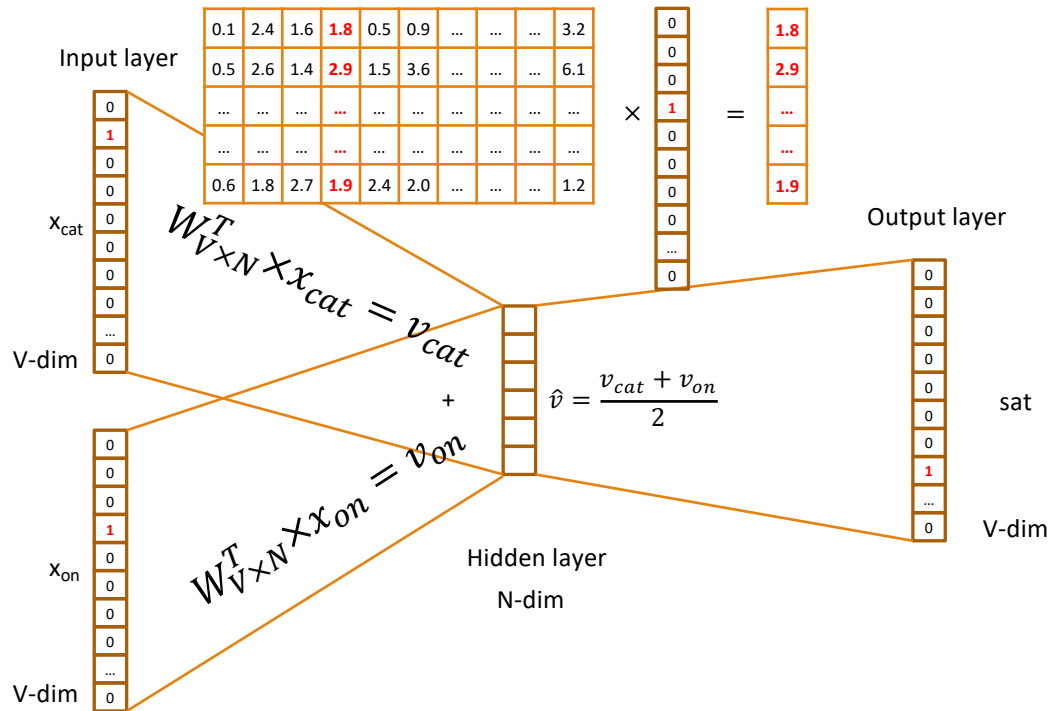
We must learn  $W$  and  $W'$

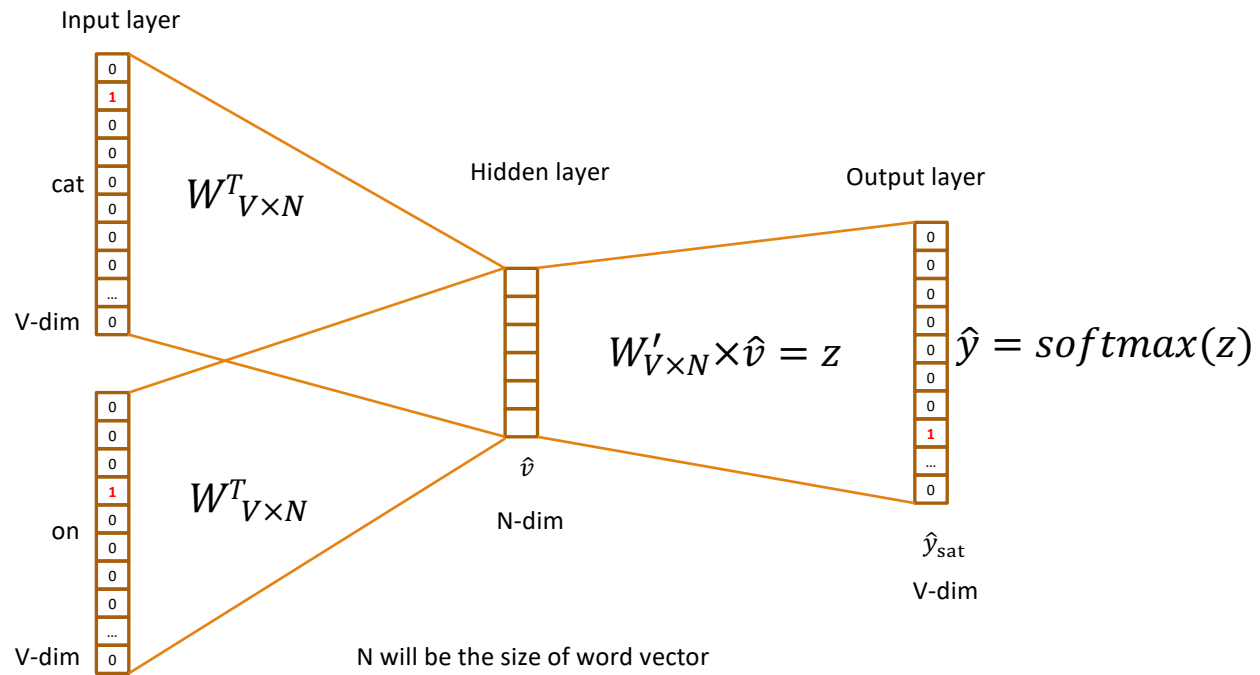



$$W_{V \times N}^T \times x_{cat} = v_{cat}$$




$$W_{V \times N}^T \times x_{on} = v_{on}$$









We can consider either  $W$  or  $W'$  as the word's representation. Or even take the average.

# Some interesting results

## Word Analogies

Test for linear relationships, examined by Mikolov et al. (2014)

a:b :: c:?



$$d = \arg \max_x \frac{(w_b - w_a + w_c)^T w_x}{||w_b - w_a + w_c||}$$

man:woman :: king:?

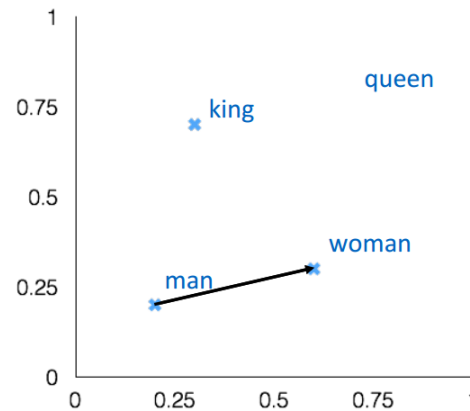
+ king [ 0.30 0.70 ]

- man [ 0.20 0.20 ]

+ woman [ 0.60 0.30 ]


---

queen [ 0.70 0.80 ]



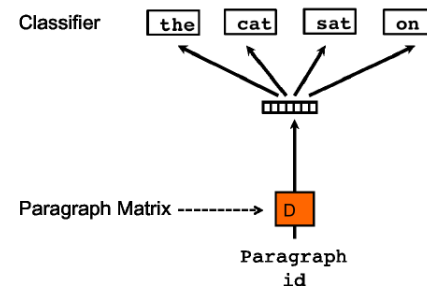
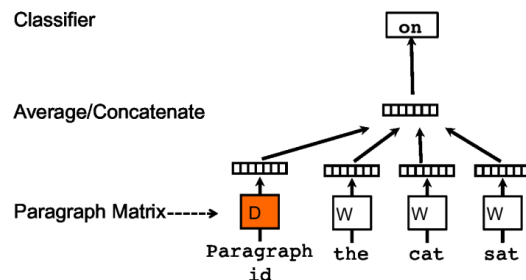


# Word analogies



# Represent the meaning of sentence/text

- Simple approach: take avg of the word2vecs of its words
- Another approach: Paragraph vector (2014, Quoc Le, Mikolov)
  - Extend word2vec to text level
  - Also two models: add paragraph vector as the input

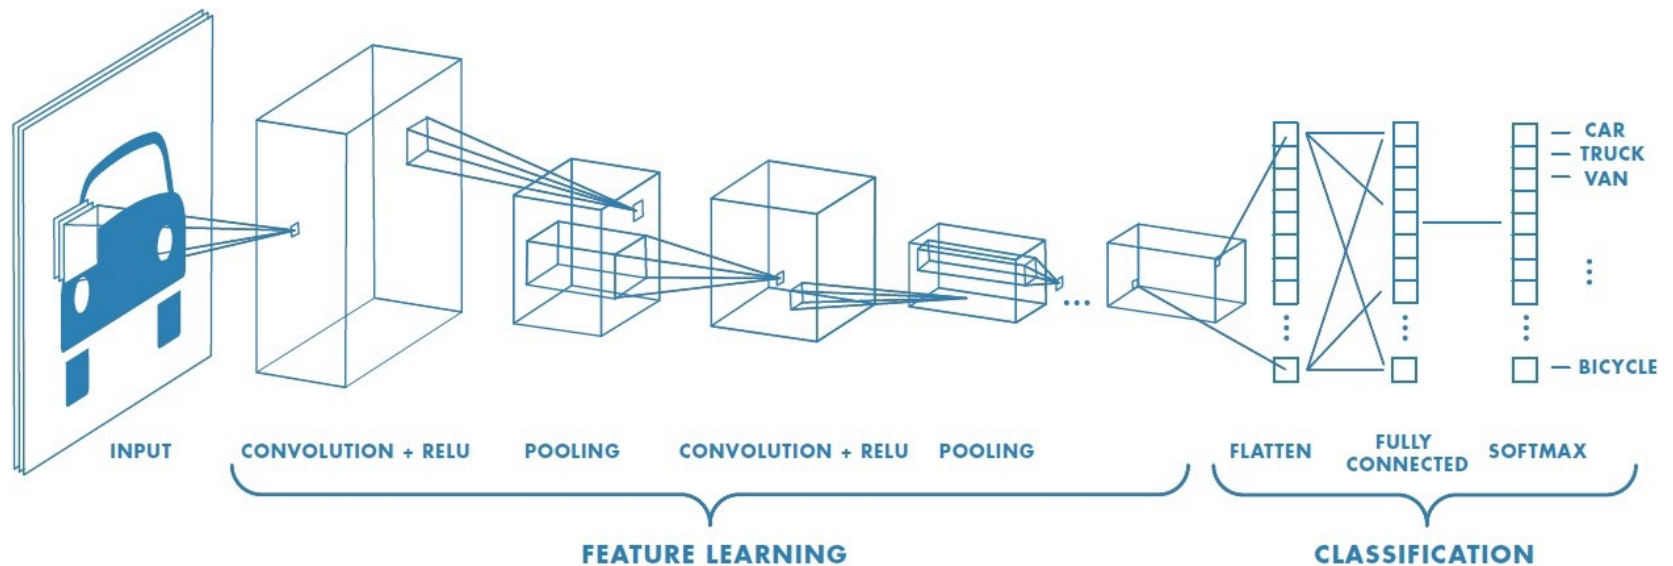


# Applications

---


- Search, e.g., query expansion
- Sentiment analysis
- Classification
- Clustering

# Recap on CNN



# CNN for Text

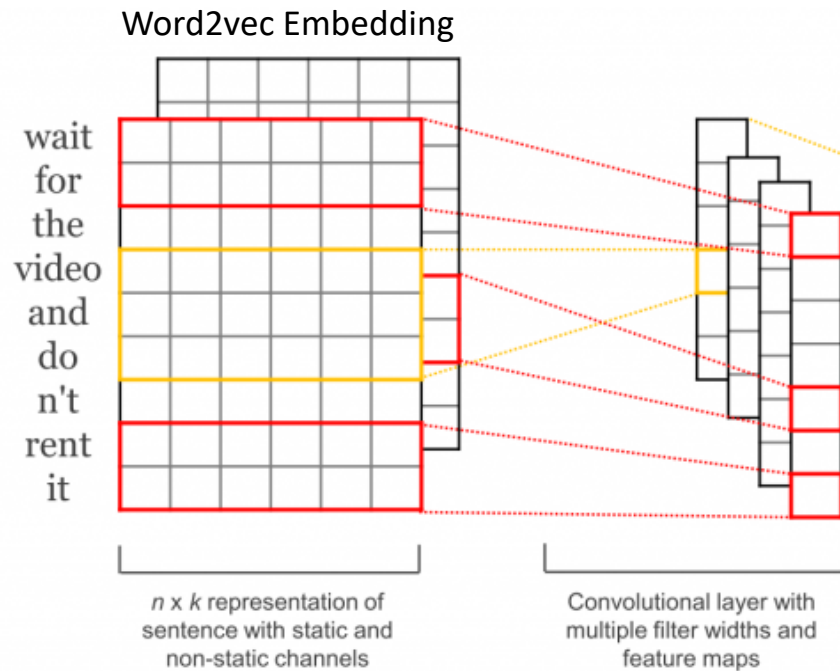
- “Shallow” CNN on sentences represented by word embeddings
  - Sentence represented as  $n \times k$  matrix ( $n$  words and embedding dimension  $k$ )



Kim, Y. (2014). Convolutional Neural Networks for Sentence Classification. In *Proceedings of the 2014 Conference on Empirical Methods in Natural Language Processing (EMNLP)*(pp. 1746-1751).

# CNN for Text

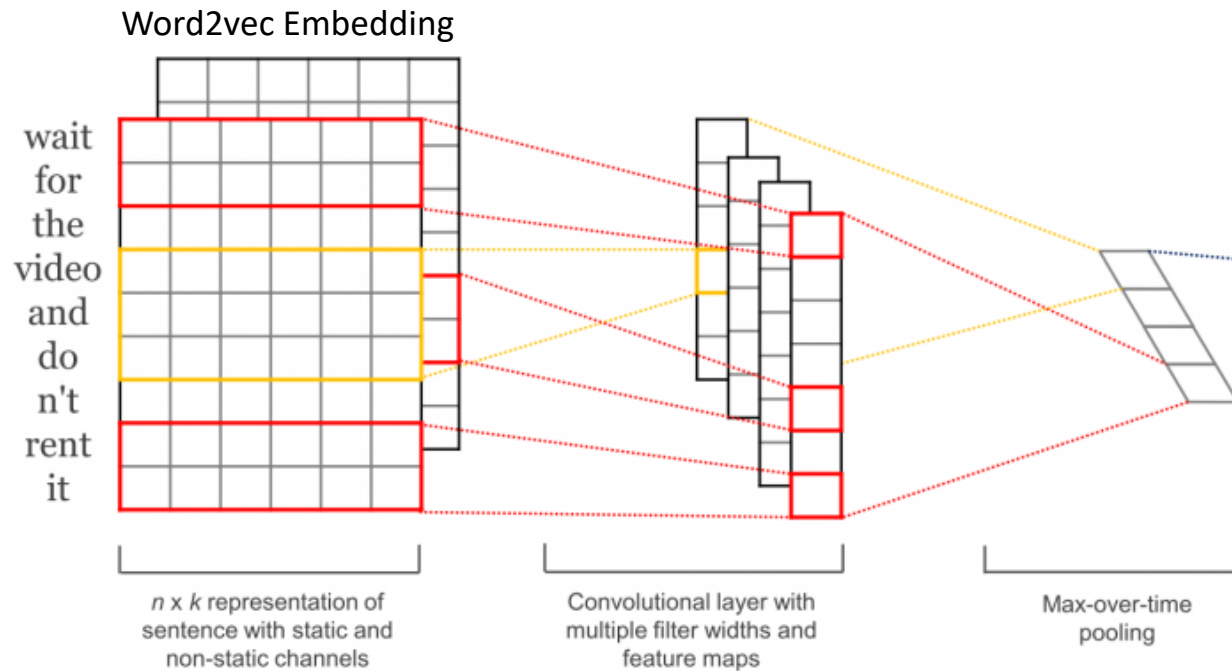
- “Shallow” CNN on sentences represented by word embeddings
  - Apply convolutional filters that cover 2,3,n words at a time (with width  $k$ )



Kim, Y. (2014). Convolutional Neural Networks for Sentence Classification. In *Proceedings of the 2014 Conference on Empirical Methods in Natural Language Processing (EMNLP)*(pp. 1746-1751).

# CNN for Text

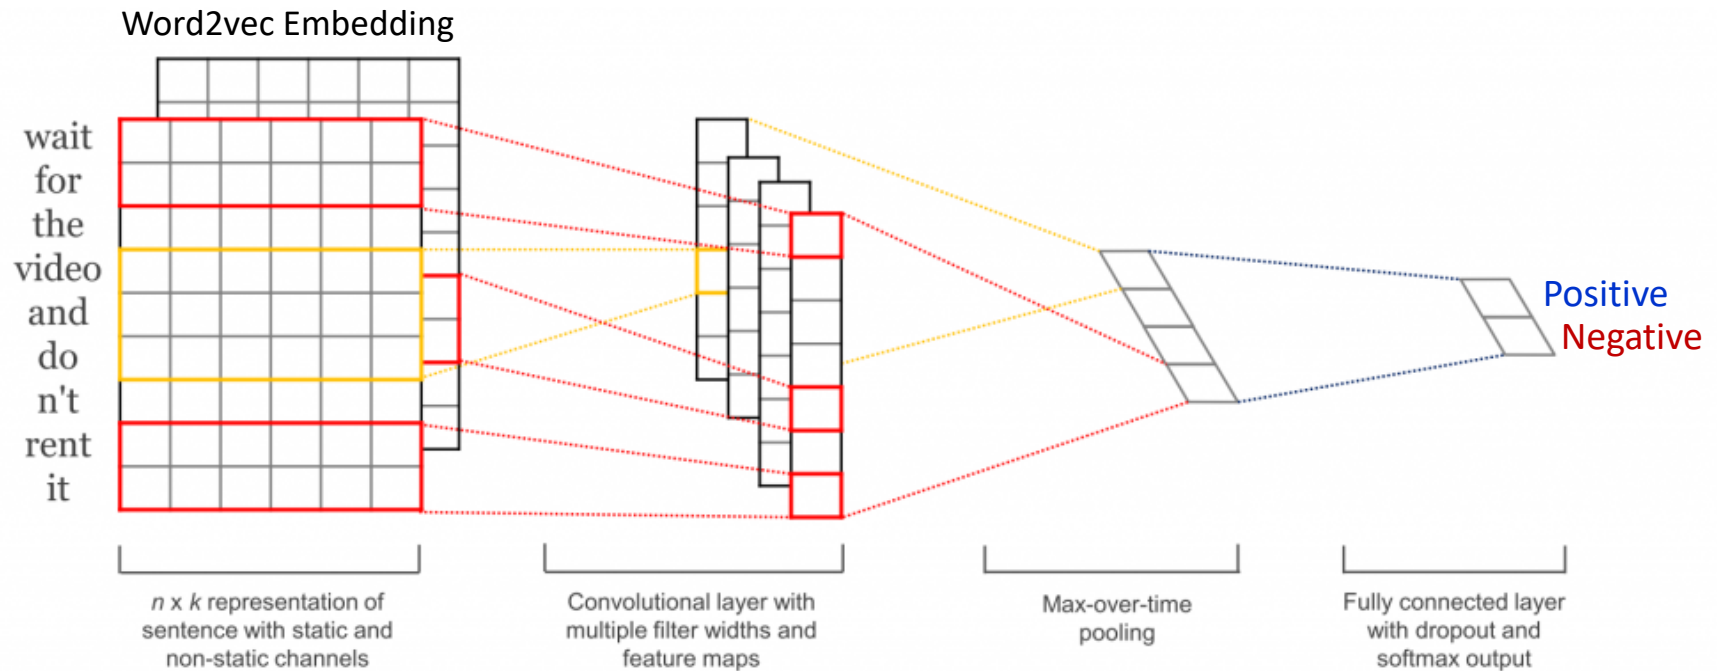
- “Shallow” CNN on sentences represented by word embeddings
  - Apply max pooling to select highest value, and flatten layer



Kim, Y. (2014). Convolutional Neural Networks for Sentence Classification. In *Proceedings of the 2014 Conference on Empirical Methods in Natural Language Processing (EMNLP)*(pp. 1746-1751).

# CNN for Text

- “Shallow” CNN on sentences represented by word embeddings
  - Final softmax layer to determine most likely class



Kim, Y. (2014). Convolutional Neural Networks for Sentence Classification. In *Proceedings of the 2014 Conference on Empirical Methods in Natural Language Processing (EMNLP)*(pp. 1746-1751).



# Exercise 1: Standard Neural Networks

---


- So far, we have covered topics on the single perceptron, MLP, and CNN.
- What is the common characteristic and limitation of these neural networks?

# Temporal Sequences

---

“This morning I took the dog for a walk.”

*Sentences*



*Heart-rates*

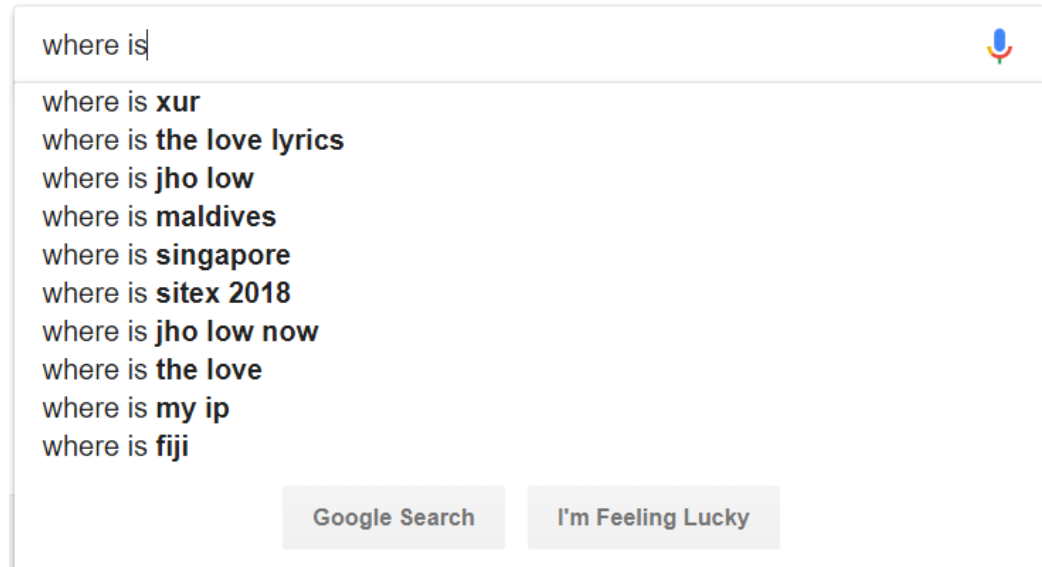


*Audio Recordings*

# Sequence Modelling Problem

---

- E.g., Sentence completion

A screenshot of a Google search interface. The search bar contains the text "where is|". Below the search bar, a list of suggestions is displayed, each starting with "where is" followed by a bolded word or phrase. The suggestions are: "xur", "the love lyrics", "jho low", "maldives", "singapore", "sitex 2018", "jho low now", "the love", "my ip", and "fiji". At the bottom of the search bar, there are two buttons: "Google Search" and "I'm Feeling Lucky".

where is|

where is **xur**  
where is **the love lyrics**  
where is **jho low**  
where is **maldives**  
where is **singapore**  
where is **sitex 2018**  
where is **jho low now**  
where is **the love**  
where is **my ip**  
where is **fiji**

Google Search I'm Feeling Lucky

# Sequence Modelling Problem

---

- E.g., Sentence completion

“This morning I took the dog for a walk.”

*given these words*

*predict what  
comes next?*

# Exercise 2: Approaches to Sequence Modelling Problem

- Using a Sentence Completion example, what are the possible problems?
  - Solution 1: Using a fixed window of preceding words

“This morning I took the dog for a walk.”

*given these 2 words, predict the next word*

- Solution 2: Model entire sentence as Bag-of-words

This morning I took the dog for a



[0 1 0 0 1 0 0 ... 0 0 1 1 0 0 0 1]

- Solution 3: Like solution 1 but with a very large window

“This morning I took the dog for a walk.”

*given these 7 words, predict the next word*

[1 0 0 0 0 0 0 0 0 0 1 0 0 0 0 0 1 0 0 0 0 1 0 0 0 0 0 0 0 0 0 1 0 0 0 ...]  
morning I took the dog ...

# Motivations behind RNN

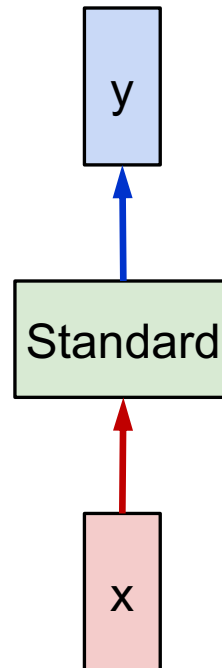
---

- Inputs and output may not be of a fixed length
  - E.g., An input paragraph of variable word count
- Want to model temporal aspect (sequence order) as context
  - E.g., language translation or stock prediction
- Hard to determine appropriate window size for context
  - E.g., past 3 days VS 6 months, previous word VS last 8<sup>th</sup> word
- Want to share parameters across sequence
  - E.g, patterns that appear in different parts of the temporal sequence

# Recurrent Neural Networks

---


- Standard Neural Networks generate a single output



# Recurrent Neural Networks

---

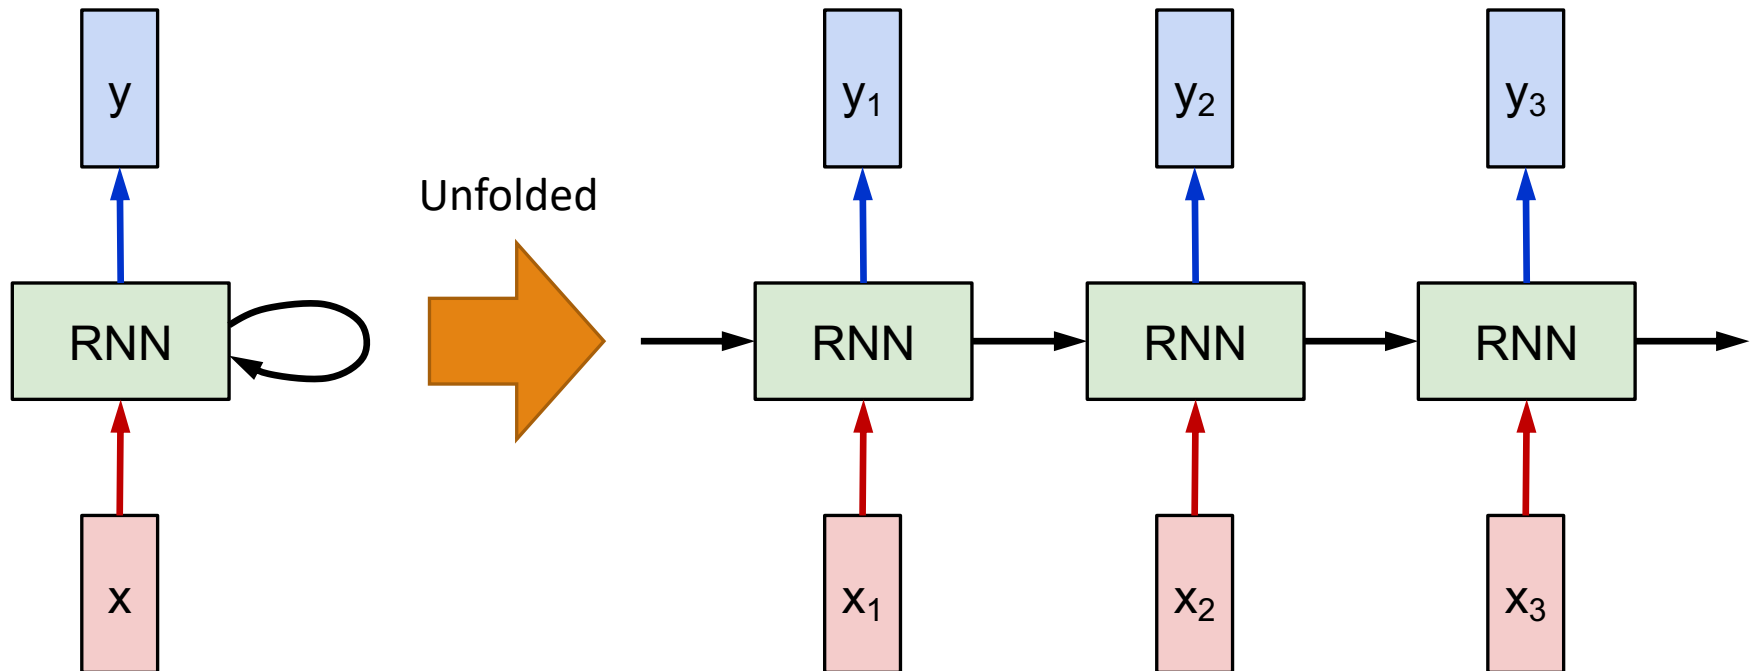
- Recurrent Neural Networks aim to process a sequence of inputs to generate a sequence of outputs at different time steps





# Recurrent Neural Networks

- Recurrent Neural Networks aim to process a sequence of inputs to generate a sequence of outputs at different time steps

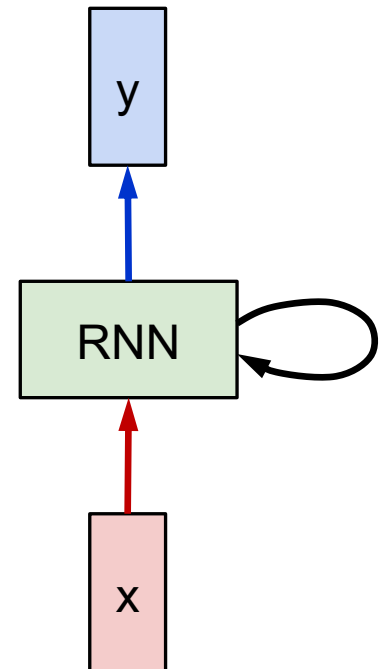


# Recurrent Neural Networks

---

- RNNs process a sequence of inputs  $X$  using a recurrence formula at each time-step  $t$

$$h_t = f_W(h_{t-1}, x_t)$$



# Recurrent Neural Networks

- RNNs process a sequence of inputs  $X$  using a recurrence formula at each time-step  $t$ 
  - Each time-step  $t$  depends on its previous time-step  $t-1$

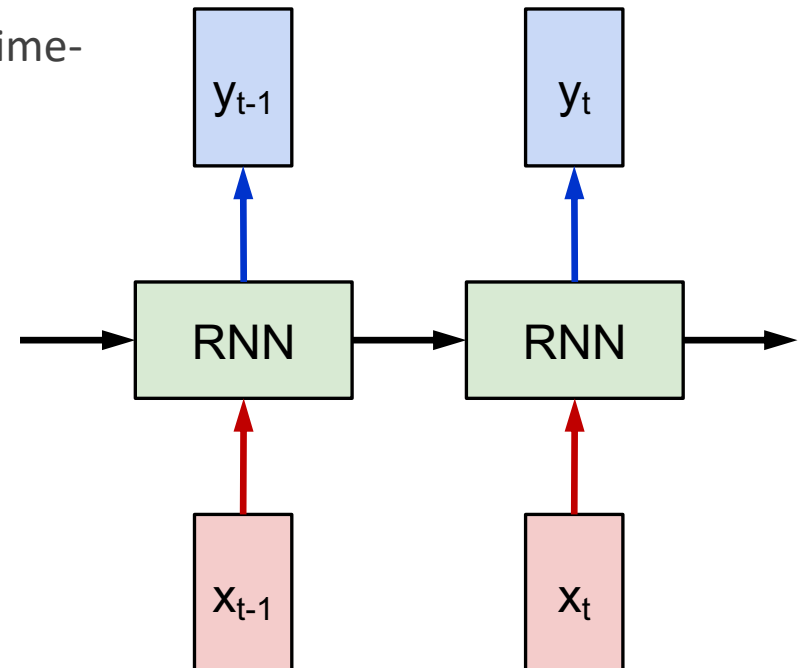
$$\boxed{h_t} = \boxed{f_W}(\boxed{h_{t-1}}, \boxed{x_t})$$

new state

some function with parameters  $W$

old state

input vector at some time step



# Recurrent Neural Networks

- RNNs process a sequence of inputs  $X$  using a recurrence formula at each time-step  $t$ 
  - The same function  $f_W$  and parameters  $W$  are shared across time-steps


$$\boxed{h_t} = \boxed{f_W}(\boxed{h_{t-1}}, \boxed{x_t})$$

new state

some function with parameters  $W$

old state

input vector at some time step



# Recurrent Neural Networks


- RNNs process a sequence of inputs  $X$  using a recurrence formula at each time-step  $t$ 
  - The same function  $f_W$  and parameters  $W$  are shared across time-steps

$$h_t = f_W(h_{t-1}, x_t)$$



$$h_t = \tanh(W_{hh}h_{t-1} + W_{xh}x_t)$$

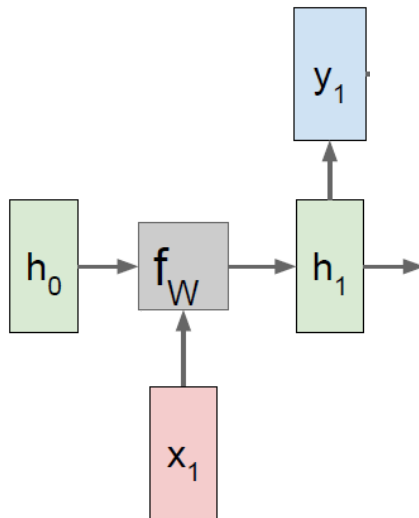
$$y_t = W_{hy}h_t$$



# RNN Example

---

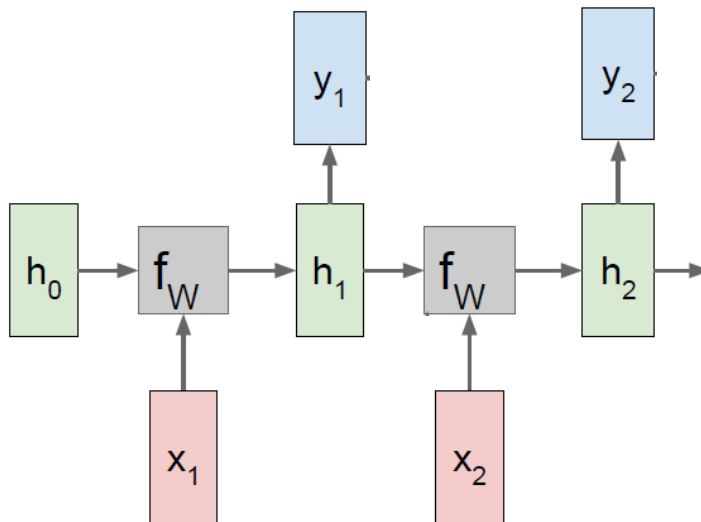
- Start from time-step 1
  - Output  $y_1$ ,



# RNN Example

---

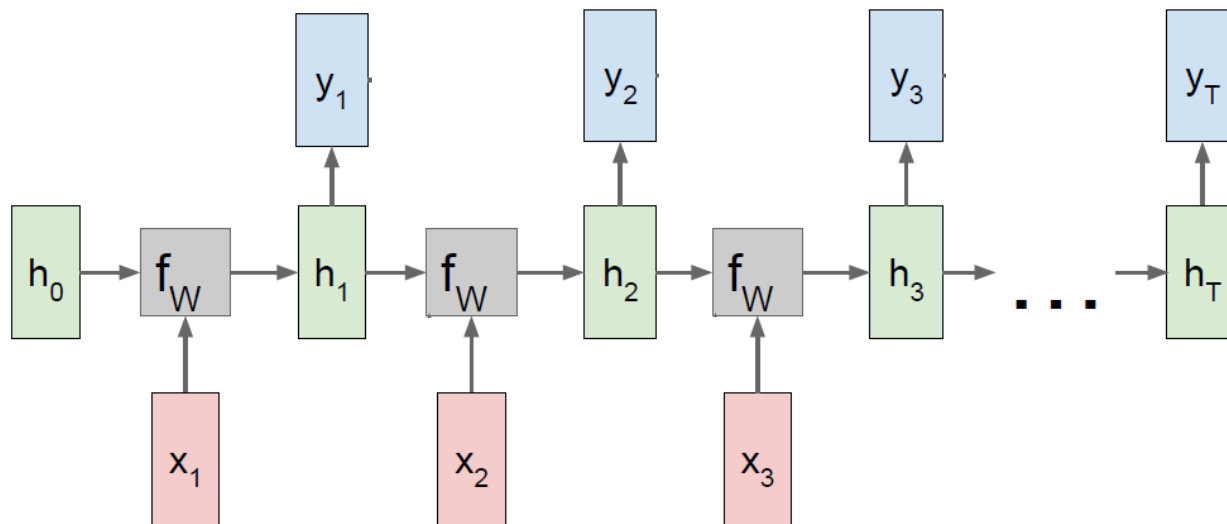
- Start from time-step 1, then time-step 2,
  - Output  $y_1, y_2$  generated



# RNN Example

---

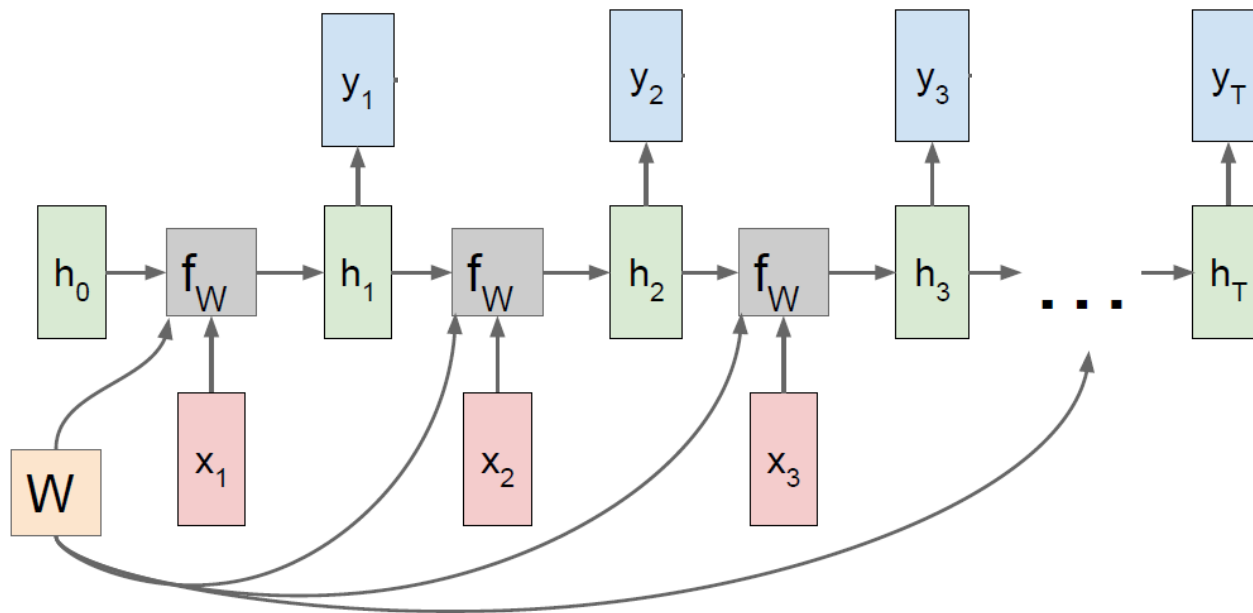
- Start from time-step 1, then time-step 2, until time-step  $T$ 
  - Output  $y_1, y_2, \dots, y_T$  generated





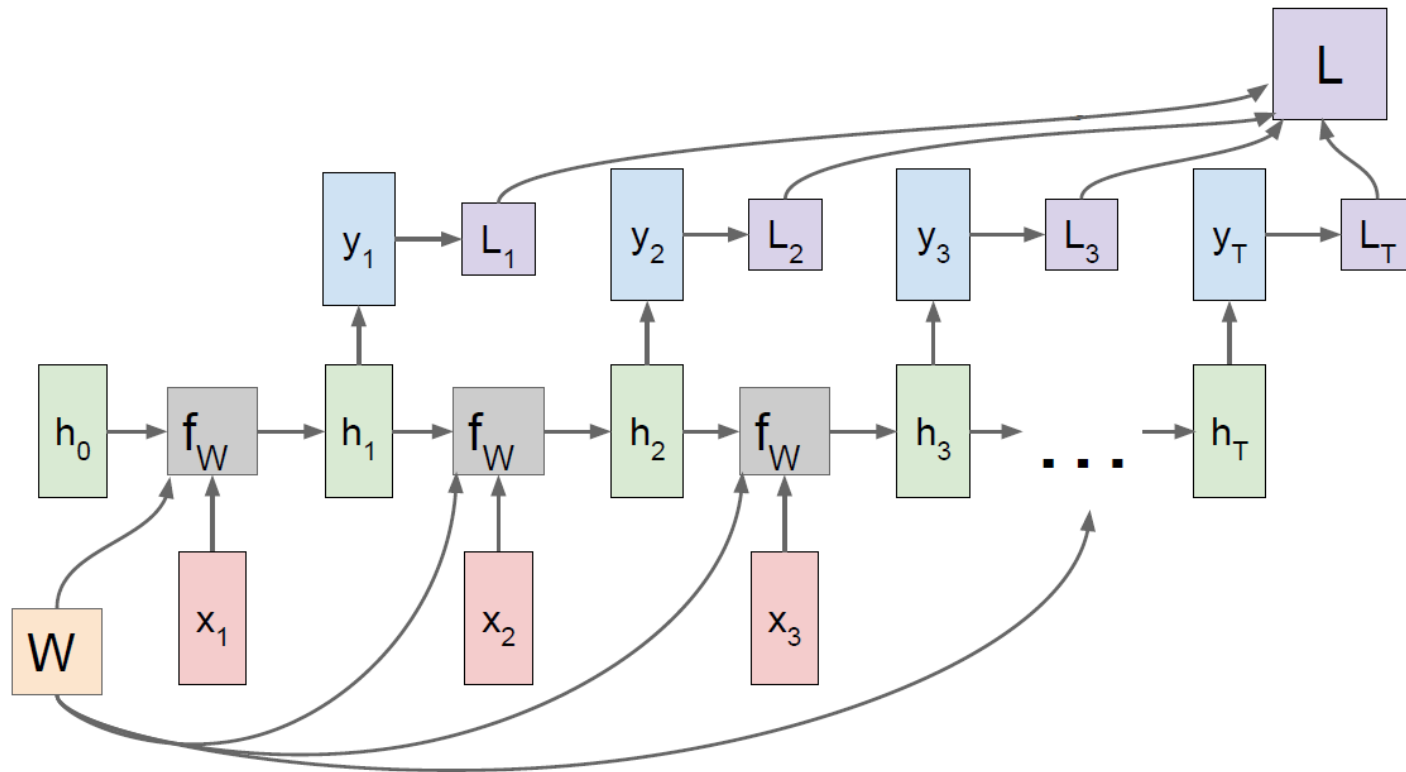
# RNN Example

- Same weights  $W$  used throughout all time-steps



# RNN Example

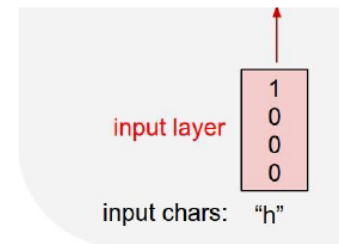
- Same weights  $W$  used throughout all time-steps
  - Learn weights by minimizing overall loss  $L$  from all time-steps



# RNN Example: Character-level

---

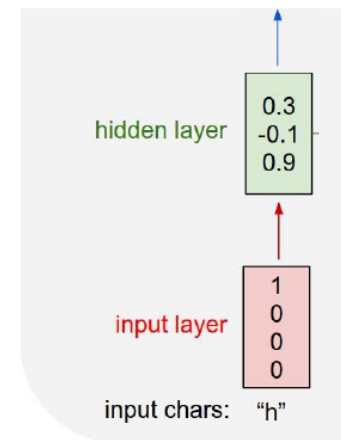
- Problem: Trying to predict a sequence of characters
  - Assuming a vocabulary of [h,e,l,o]



# RNN Example: Character-level

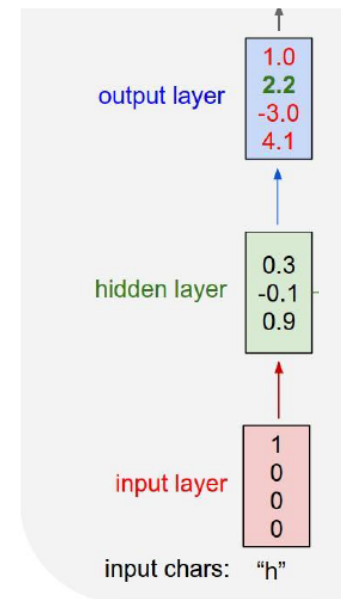
---

- Problem: Trying to predict a sequence of characters
  - Assuming a vocabulary of [h,e,l,o]



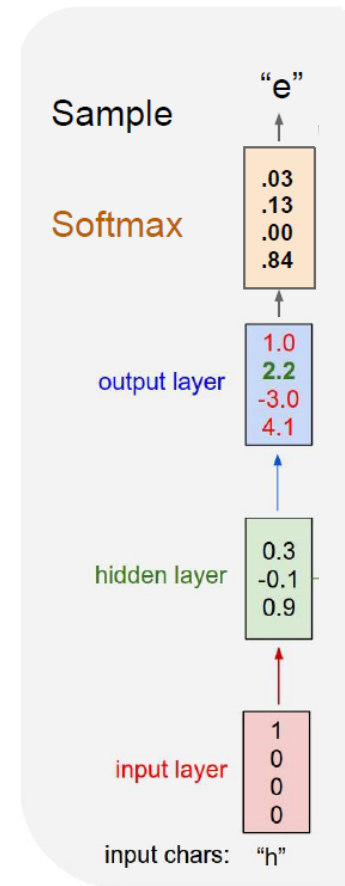
# RNN Example: Character-level

- Problem: Trying to predict a sequence of characters
  - Assuming a vocabulary of [h,e,l,o]



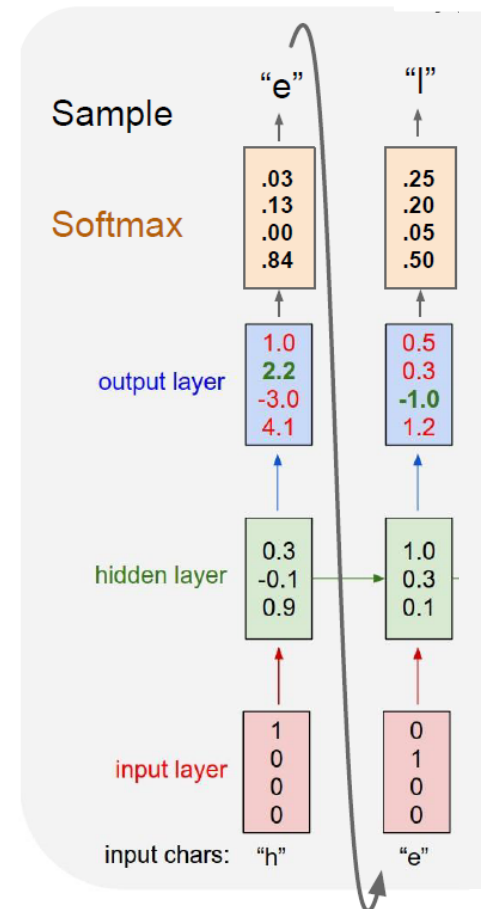
# RNN Example: Character-level

- Problem: Trying to predict a sequence of characters
  - Assuming a vocabulary of [h,e,l,o]



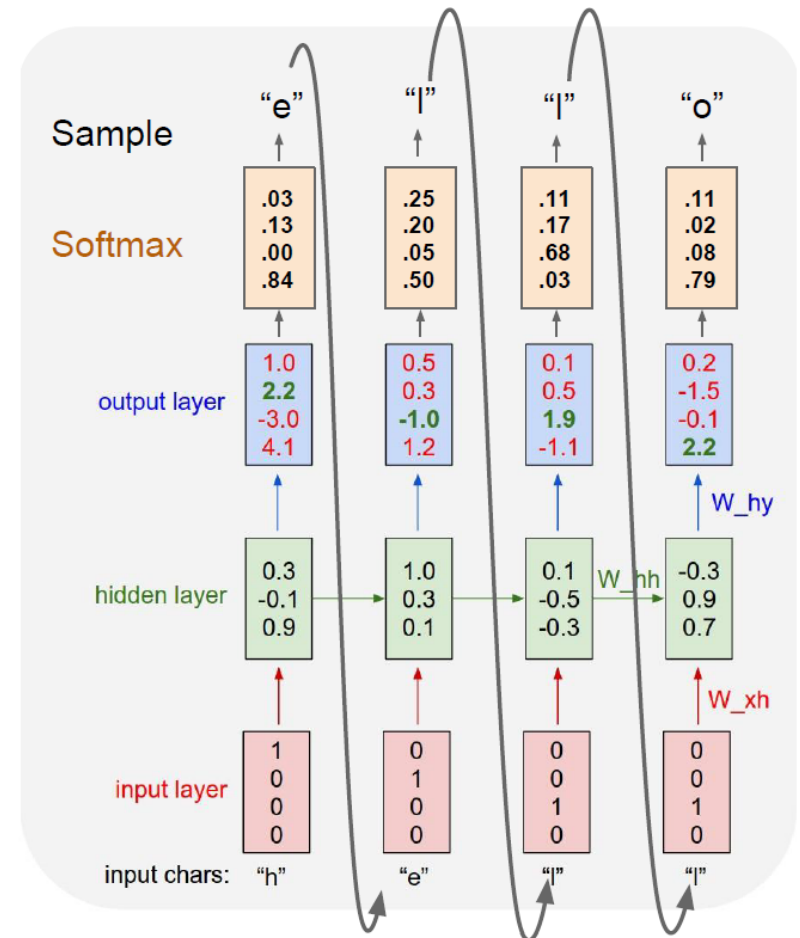
# RNN Example: Character-level

- Problem: Trying to predict a sequence of characters
  - Assuming a vocabulary of [h,e,l,o]



# RNN Example: Character-level

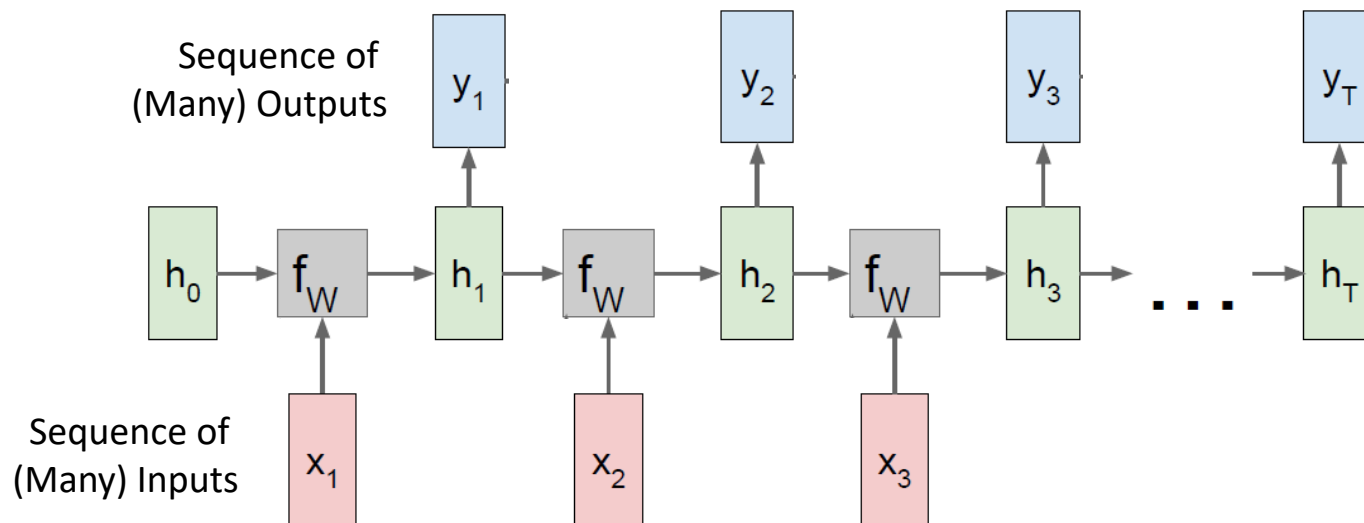
- Problem: Trying to predict a sequence of characters
  - Assuming a vocabulary of [h,e,l,o]





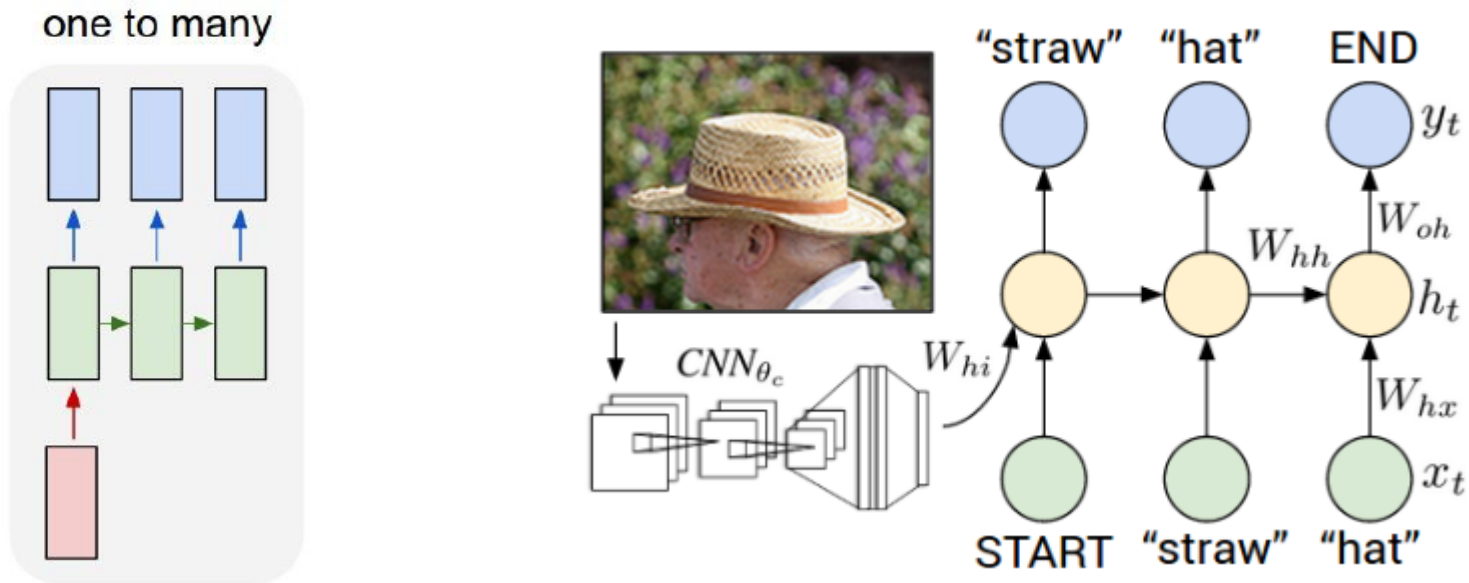
# Exercise 3: Types of RNN

- Based on the input/output sequence, we have previously examined a many-to-many type of RNN. What other types can you think of?



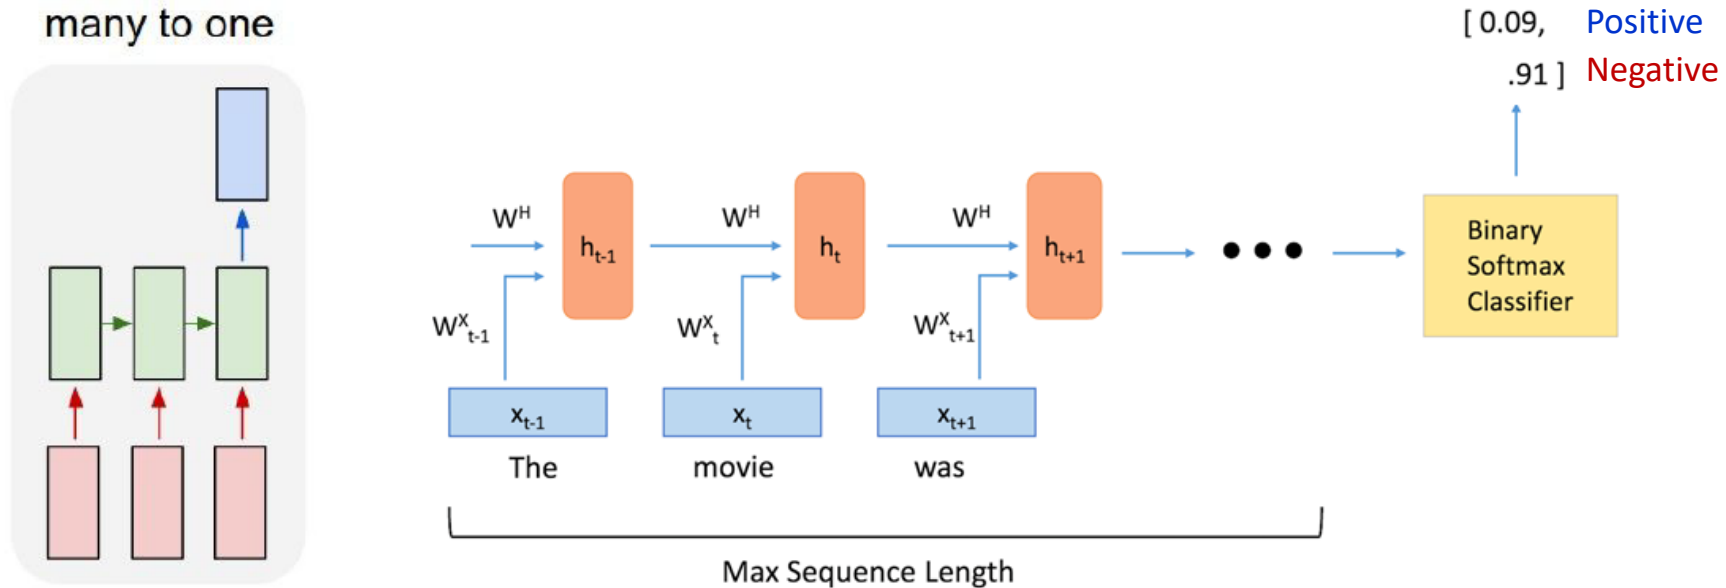
# RNN Applications

- One-to-many: Image Captioning



# RNN Applications

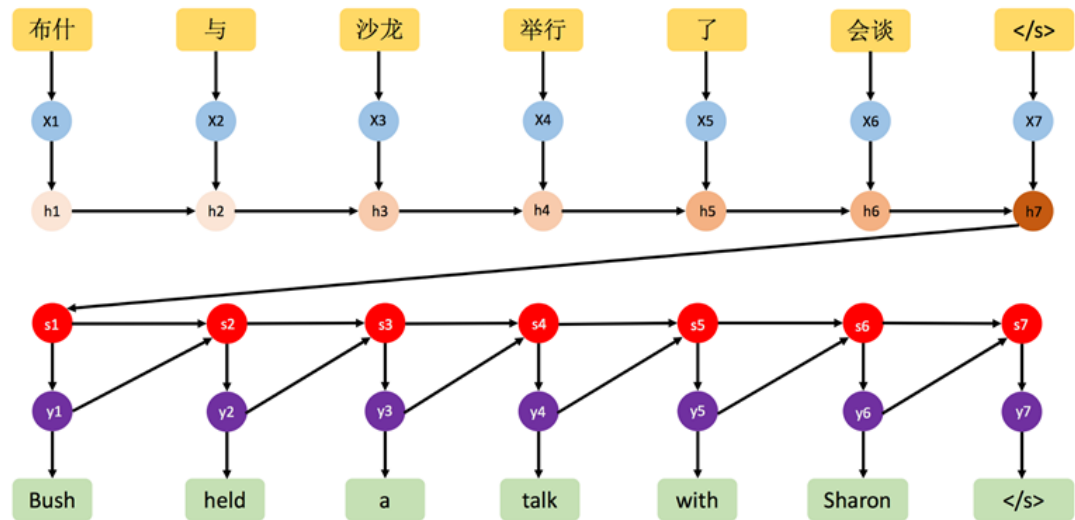
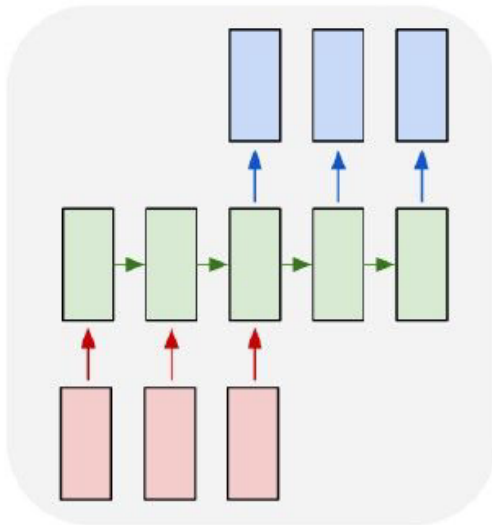
- Many-to-one: Sentiment Classification



# RNN Applications

- Many-to-many: Language Translation

many to many

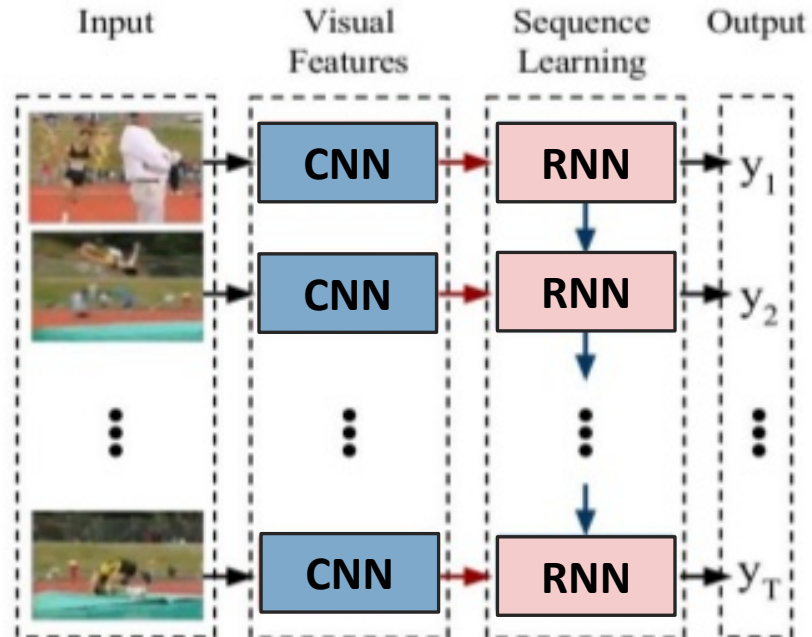



(Sutskever et al., 2014)

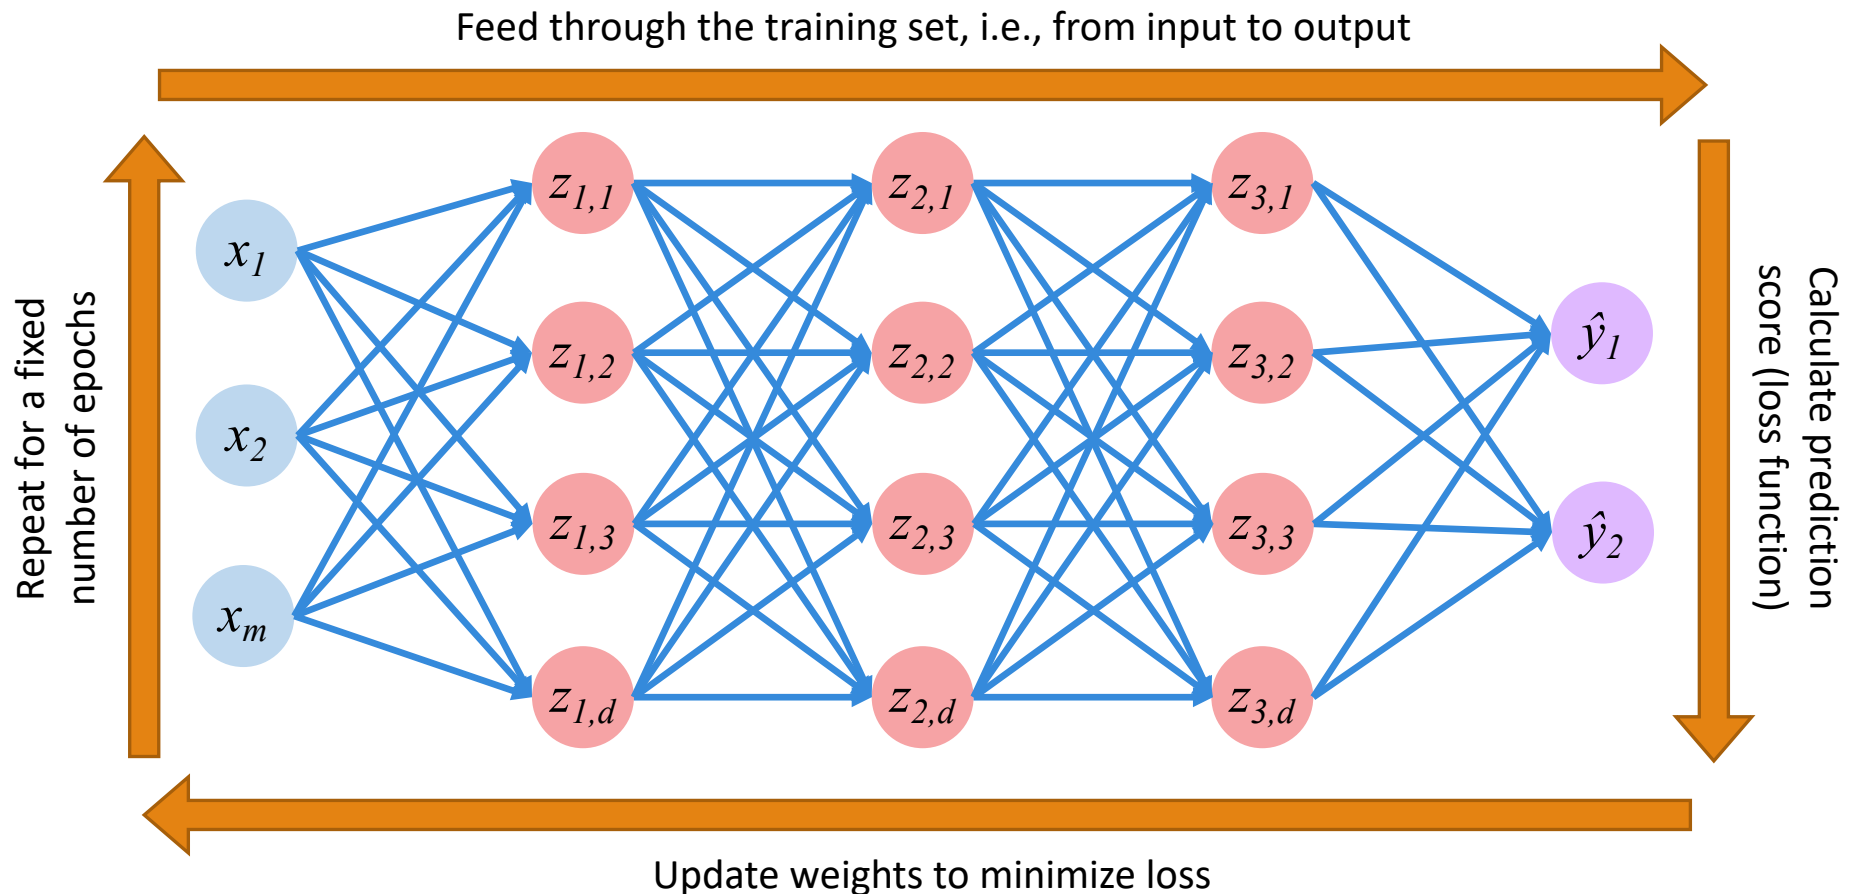
# RNN Applications

- Many-to-many: Video Classification

many to many

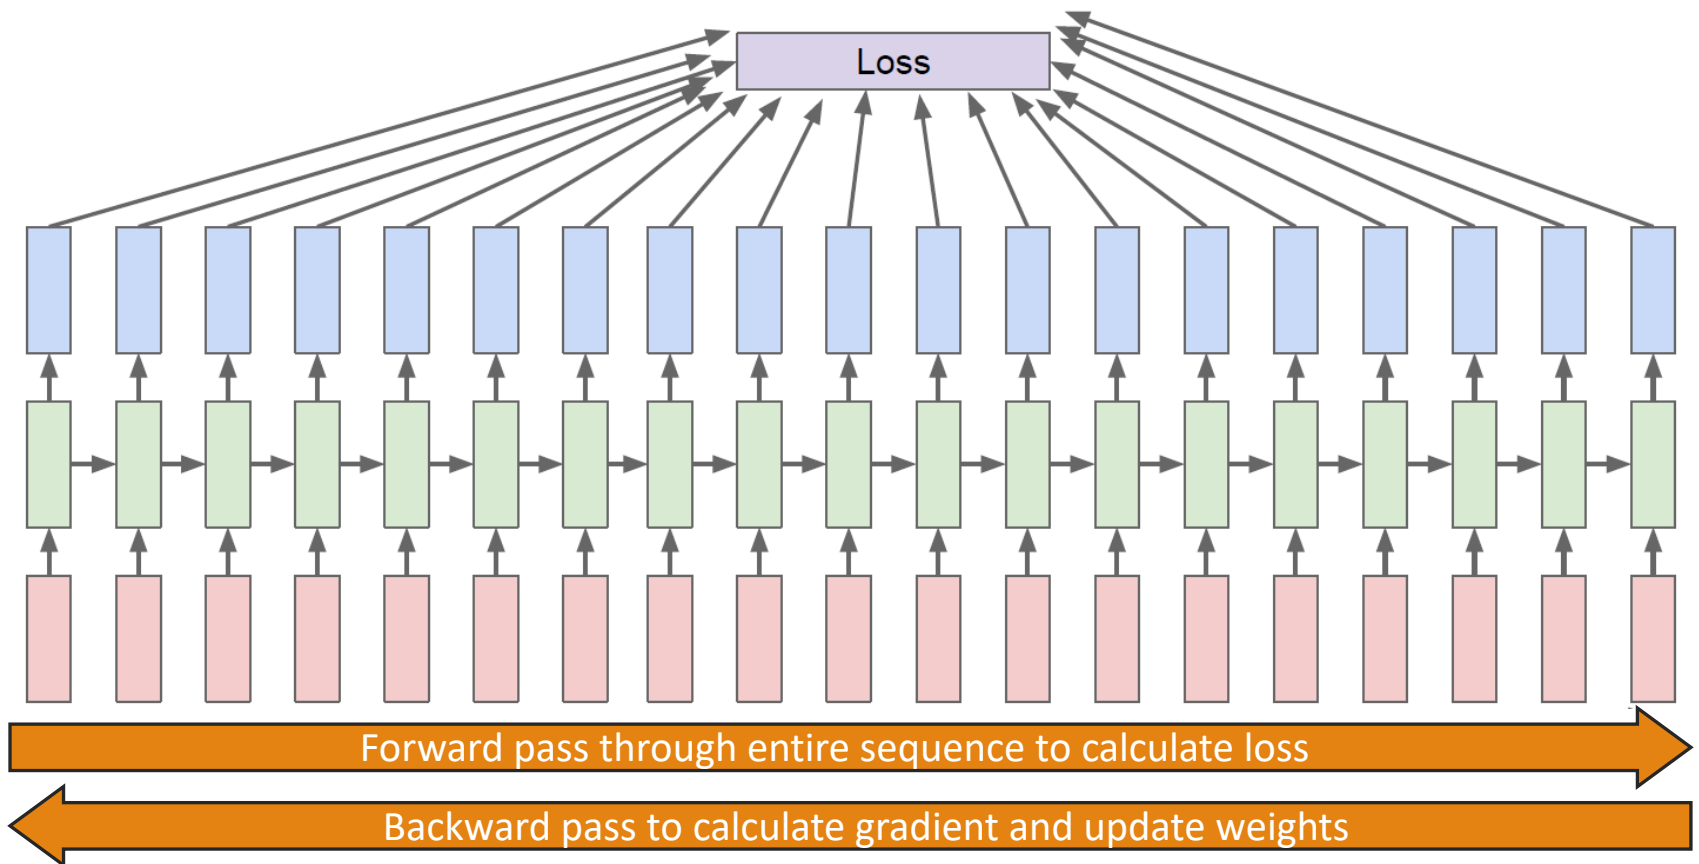


# Recap on Training MLP



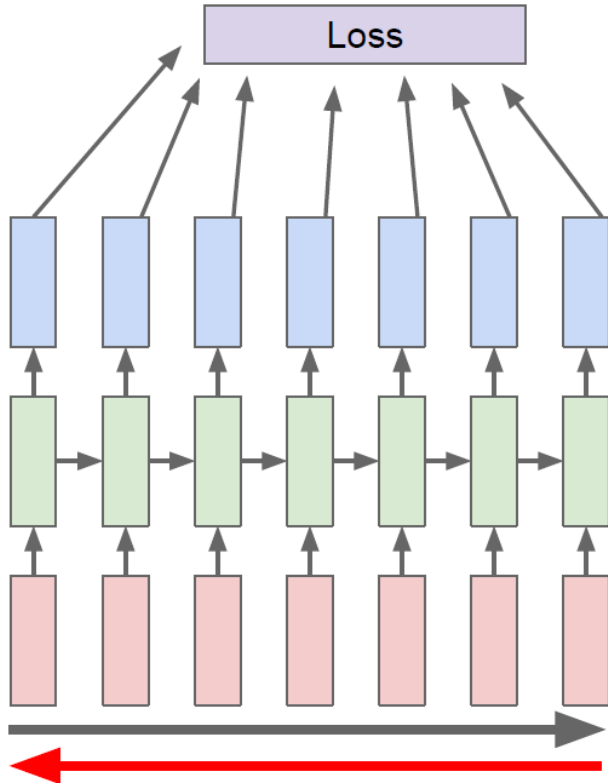
# RNN Training

- Backpropagation through time



# RNN Training

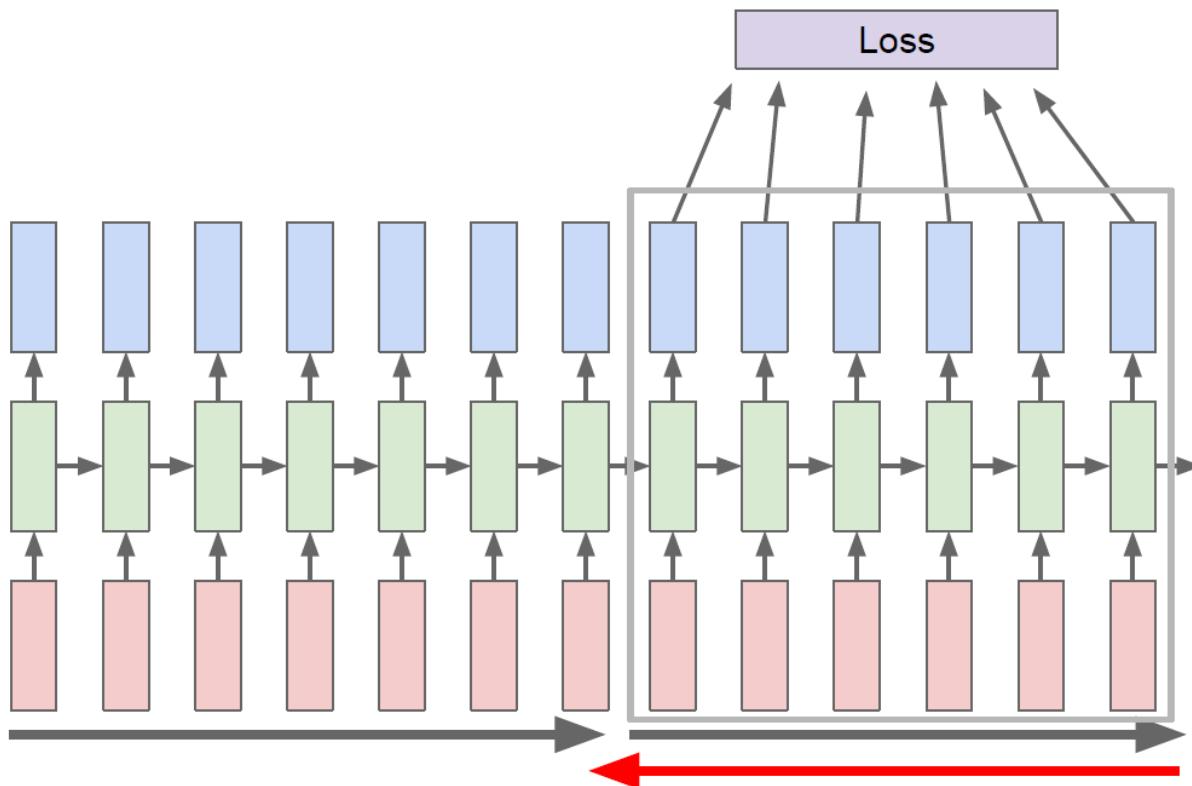
- Truncated Backpropagation through time
  - Instead of the entire sequence, break it up into smaller sub-sequences






# RNN Training

- Truncated Backpropagation through time
  - Perform the forward/backward pass for each sub-sequence



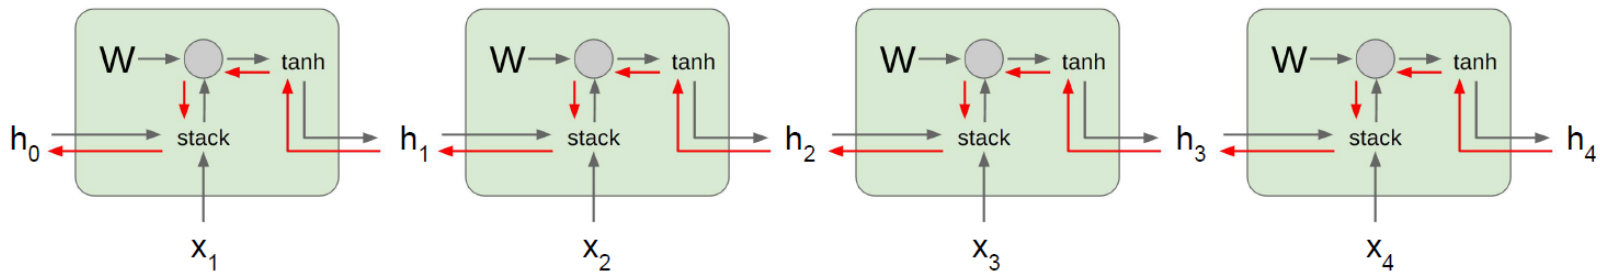
# RNN Training

- Truncated Backpropagation through time



# Problems with Vanilla RNNs

- Vanishing and Exploding Gradients



- Unable to remember inputs from long ago

*I live in **France** ... .. I speak fluent **French**.*

